



# HTTP: Dual-Stack TCP Server

Prepared by **Hamza Ghaffar** (still in progress for more complex topics to learn)

## Header Types

### 1 Request Headers

**Host:** Specifies the domain name of the server (useful when multiple domains are hosted on a single IP). [www.example.com](http://www.example.com)

**User-Agent:** Contains information about the client software (browser type, version, operating system)

**Accept:** Specifies the media types that are acceptable for the response (e.g., HTML, JSON).

Example: **Accept: text/html,application/xhtml+xml**

**Accept-Language:** Indicates the preferred language for the response.

- Example: **Accept-Language: en-US,en;q=0.5**

**Connection:** Indicates whether the client wants to keep the connection alive after the current request.

- Example: **Connection: keep-alive**

**Cache-Control:** Directives for caching mechanisms in both requests and responses.

- Example: **Cache-Control: no-cache**

### 2 Response Headers

**Status:** Indicates the result of the request, including the status code and a description.

- Example: **HTTP/1.1 200 OK**

**Content-Type:** Indicates the media type of the resource being sent.

- Example: **Content-Type: application/json**

**Content-Length:** The size of the response body in bytes.

- Example: **Content-Length: 1234**

**Server:** Contains information about the server software.

- Example: **Server: Apache/2.4.1 (Unix)**

**Cache-Control:** Controls how the response should be cached.

- Example: **Cache-Control: no-store**

**Set-Cookie:** Used to send cookies from the server to the client.

- Example: **Set-Cookie: sessionId=abc123; HttpOnly**



## 1.1 Connection creation

```
dual_stack_socket.bind((host, port))
```

```
dual_stack_socket.listen()
```

```
"""the accept() function is a blocking call. This means that when
dual_stack_socket.accept() is reached, the server "pauses" here until a client
connects."""
client_socket, address = dual_stack_socket.accept()
```

## 1.2 Receiving the Request

```
request = client_socket.recv(1024).decode()

"""
# Request Line
GET /index.html HTTP/1.1
"""

# Headers - Key: Value

Host: example.com
User-Agent: Mozilla/5.0
Accept-Language: en-US,en;q=0.5
Connection: keep-alive
"""

"""
# Body
{
    "key": "value"
}
"""

lines = request.splitlines()
if len(lines) > 0:
    request_line = lines[0]
    method, path, _ = request_line.split()
```

## 2.1 Preparing the Response

```

"""
Prepare the HTTP response headers.
Status Line: This line indicates the HTTP version, the status code, and a status
message

Response Headers: These headers provide additional information about the
response, similar to request headers.

Blank-Space:

Response Body: This is the actual content being sent back to the client, such as
HTML, JSON, images, etc.

"""

response = f"HTTP/1.1 200 OK\r\n" \
           f"Content-Type: application/json\r\n" \
           f"Content-Length: {len(response_body)}\r\n" \
           f"\r\n" \
           f"{response_body}"

# Send the response
client_socket.sendall(response.encode())

```

This guide will include a proper control flow to help you understand how everything fits together.

- **app/main.py:** This is the core file where your HTTP server code resides. It's the main entry point for your application's logic.
- **your\_program.sh:** This is a shell script that provides a simple way to run your application locally.
- **.codecrafters/run.sh:** A specialized script used for the CodeCrafters platform, which provides an environment for running and testing code (such as server simulations).
- **Pipfile:** Defines your Python dependencies, listing the libraries required to run your application.
- **Pipfile.lock:** Locks the exact versions of dependencies used in the project to ensure consistency across different environments.

### your\_program.sh

```

#!/bin/bash

set -e # Exit immediately if a command exits with a non-zero status

# Check if pipenv environment is set up and install dependencies if not
if [ ! -d ".venv" ]; then
    echo "Setting up the environment with pipenv..."
    pipenv install # Install dependencies and create a virtual environment
fi

# Start the HTTP server within the pipenv environment
echo "Starting the HTTP server..."

```

```
exec pipenv run python3 -m app.main "$@"
```



### Pipfile

The **Pipfile** is a configuration file used by Pipenv, a dependency manager for Python projects. This file defines the dependencies required for your project[While this file does not directly contain server code, it's crucial for ensuring you have the right packages installed to support your development.]



### Pipfile.lock

Similar to **Pipfile**, this file locks the versions of the dependencies. It ensures that the same versions are used every time the project is set up, maintaining consistency across different environments. Like **Pipfile**, it's essential for the server but does not contain server code.

## Improved Header Handling

```
# Step 1: Improved Header Handling
# Extracting All Headers: Implement a method to read and parse all request
headers to make decisions based on them.

headers = {}
for line in lines[1:]:
    if ':' in line:
        key, value = line.split(':', 1)
        headers[key] = value

# Display the headers in the console
print("Extracted Headers:")
for key, value in headers.items():
    print(f"{key}: {value}")
```

## Read the Request Body

This is important because only **POST** requests will have a body that you want to read. If it's not a POST request, the server does not need to look for a body.

By default, in this approach, we cannot handle concurrent connections for more than 1 client.  
~~its old fashion~~

```
# Read the Request Body
# Handling POST Requests: Implement logic to read the request body for POST
requests.
if method == "POST":
    content_length = int(headers.get('Content-Length', 0))
    body = request[request.index("\r\n\r\n") + 4:][:content_length]
    logging.info(f"Received body: {body}") # Log the received body
```

Complete - example code

```

import socket
import logging
import json

logging.basicConfig(level=logging.INFO, format='%(asctime)s - %(levelname)s - %(message)s')

def start_server(host='::', port=4221):
    try:
        dual_stack_socket = socket.socket(socket.AF_INET6, socket.SOCK_STREAM)
        dual_stack_socket.setsockopt(socket.IPPROTO_IPV6, socket.IPV6_V6ONLY, 0)

        dual_stack_socket.bind((host, port))
        logging.info(f"Server listening on {host}:{port}...")

        dual_stack_socket.listen()

    while True:
        client_socket, address = dual_stack_socket.accept()
        logging.info(f"Connection from {address} has been established!")
        logging.info("Client socket info:")
        logging.info(f"Client Address: {address}")
        logging.info(f"Socket Family: {client_socket.family}")
        logging.info(f"Socket Type: {client_socket.type}")
        logging.info(f"Protocol: {client_socket.proto}")
        logging.info(f"File Descriptor: {client_socket.fileno()}")

        try:
            while True:
                request = client_socket.recv(1024).decode()
                if not request:
                    break

                logging.info("Received request:")
                logging.info(request)

                lines = request.splitlines()
                if len(lines) > 0:
                    request_line = lines[0]
                    method, path, _ = request_line.split()

                    # Step 1: Improved Header Handling
                    # Extracting All Headers: Implement a method to read and
                    # parse all request headers to make decisions based on them.
                    headers = {}
                    for line in lines[1:]:
                        if ': ' in line:
                            key, value = line.split(': ', 1)
                            headers[key] = value

                    # Display the headers in the console
                    print("Extracted Headers:")
                    for key, value in headers.items():
                        print(f"{key}: {value}")

                    # Step 2: Read the Request Body
                    # Handling POST Requests: Implement logic to read the
                    # request body for POST requests.
                    if method == "POST":
                        content_length = int(headers.get('Content-Length',
0))

                        body = request[request.index("\r\n\r\n") +

```

Now we can achieve this setup via `asyncio`

```

import asyncio
import logging
import json

logging.basicConfig(level=logging.INFO, format='%(asctime)s - %(levelname)s - %(message)s')

async def handle_client(reader, writer):
    address = writer.get_extra_info('peername')
    logging.info(f"Connection from {address} has been established!")

    try:
        while True:
            data = await reader.read(1024)
            if not data:
                break # Client closed the connection

            request = data.decode()
            logging.info("Received request:")
            logging.info(request)

            # Process the request
            lines = request.splitlines()
            if len(lines) > 0:
                request_line = lines[0]
                method, path, _ = request_line.split()

                # Step 1: Improved Header Handling
                headers = {}
                for line in lines[1:]:
                    if ':' in line:
                        key, value = line.split(':', 1)
                        headers[key] = value

                logging.info("Extracted Headers:")
                for key, value in headers.items():
                    logging.info(f"{key}: {value}")

                # Step 2: Read the Request Body
                if method == "POST":
                    content_length = int(headers.get('Content-Length', 0))
                    body = request[request.index("\r\n\r\n") +
4:][:content_length]
                    logging.info(f"Received body: {body}")

                # Prepare the response
                response_content = {
                    "message": "Hamza Ghaffar | Request received Successfully!",
                    "method": method,
                    "path": path
                }
                response_body = json.dumps(response_content)

                response = f"HTTP/1.1 200 OK\r\n" \
                    f"Content-Type: application/json\r\n" \
                    f"Content-Length: {len(response_body)}\r\n" \
                    f"\r\n" \
                    f"{response_body}"

```

**Return a File -Example**



```

import asyncio
import logging
import json
import os

# Configure logging
logging.basicConfig(level=logging.INFO, format='%(asctime)s - %(levelname)s - %(message)s')

# Use absolute path for the file
file_path = os.path.join(os.path.dirname(__file__), 'your_file.txt') # Ensure correct path

async def handle_client(reader, writer):
    address = writer.get_extra_info('peername')
    logging.info(f"Connection from {address} has been established!")

    try:
        while True:
            data = await reader.read(1024)
            if not data:
                break # Client closed the connection

            request = data.decode()
            logging.info("Received request:")
            logging.info(request)

            # Process the request
            lines = request.splitlines()
            if len(lines) > 0:
                request_line = lines[0]
                method, path, _ = request_line.split()

                # Improved Header Handling
                headers = {}
                for line in lines[1:]:
                    if ': ' in line:
                        key, value = line.split(': ', 1)
                        headers[key] = value

                logging.info("Extracted Headers:")
                for key, value in headers.items():
                    logging.info(f"{key}: {value}")

                # Read the Request Body if POST
                if method == "POST":
                    content_length = int(headers.get('Content-Length', 0))
                    body = request[request.index("\r\n\r\n") +
4:][:content_length]
                    logging.info(f"Received body: {body}")

                # Prepare the response based on method and path
                if method == "GET" and path == "/your_file": # Change this path as needed

                    try:
                        with open(file_path, 'r') as file:
                            file_contents = file.read()


                        response_body = file_contents
                        response = f"HTTP/1.1 200 OK\r\n" \
                                f"Content-Type: text/plain\r\n" \
                                f"Content-Length: {len(response_body)}\r\n" \

```



You could use third-party libraries like Twisted or asyncio for more advanced asynchronous networking capabilities

### Advanced Asynchronous Networking



Custom Loggers | Multiple Handlers | Logging Exceptions | Thread Safety



Grafana and Prometheus



General headers, entity headers, and custom headers that serve different purposes in the communication process.



**Enhanced Logging:** Implement structured logging using libraries like **structlog** to track requests and responses in production.



Security and Validation



Use a Reverse Proxy



Containerization